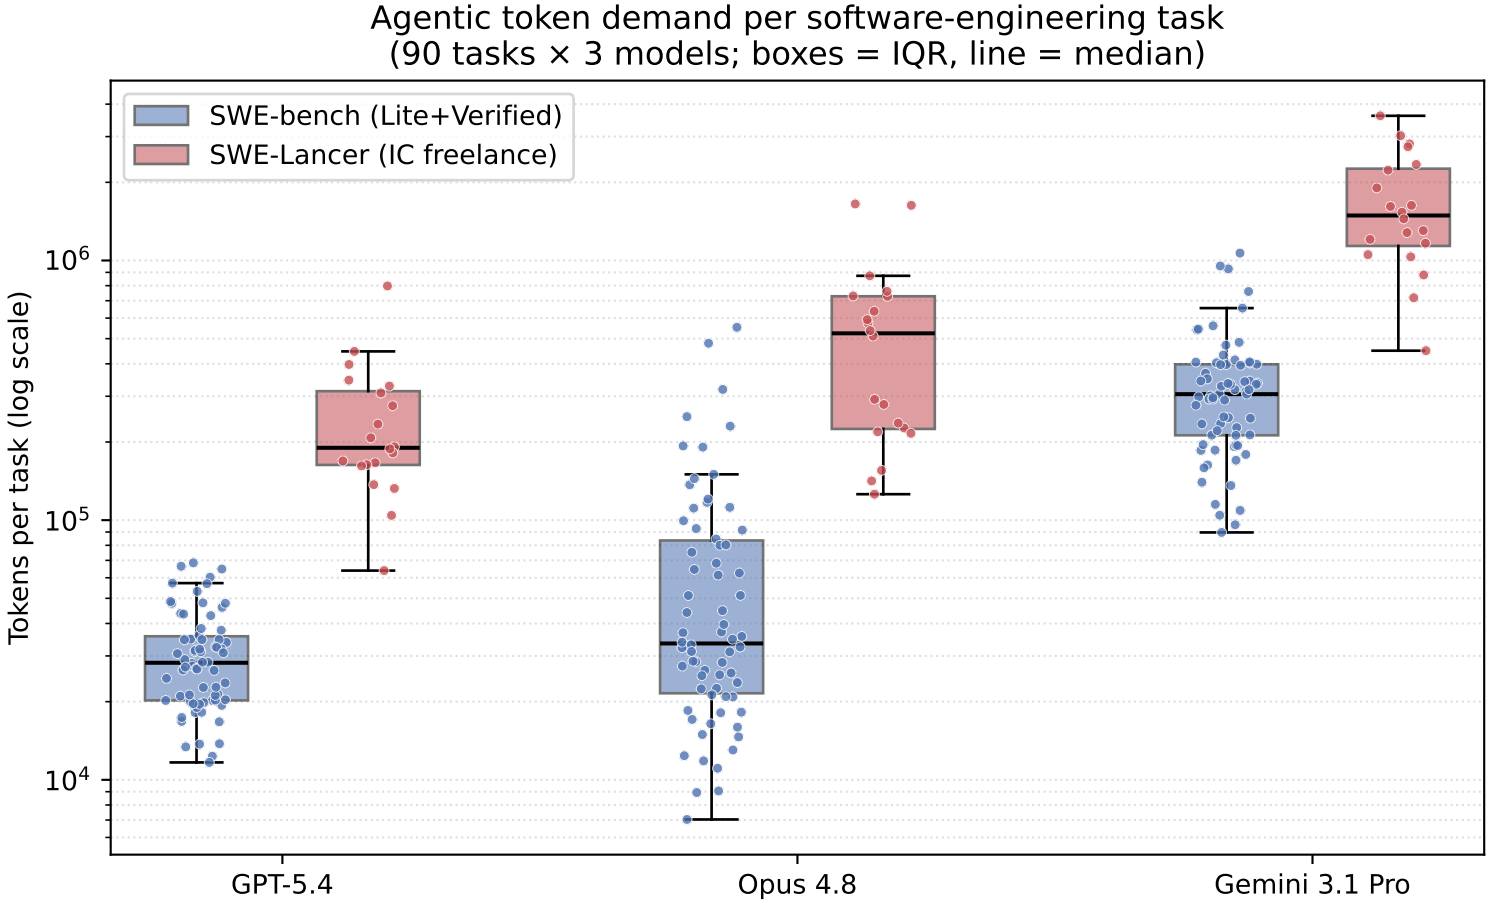
Agentic token demand per software-engineering task
(90 tasks × 3 models; boxes = IQR, line = median)
SWE-bench (Lite+Verified)
SWE-Lancer (IC freelance)
Tokens per task (log scale)
10^6
10^5
10^4
GPT-5.4
Opus 4.8
Gemini 3.1 Pro