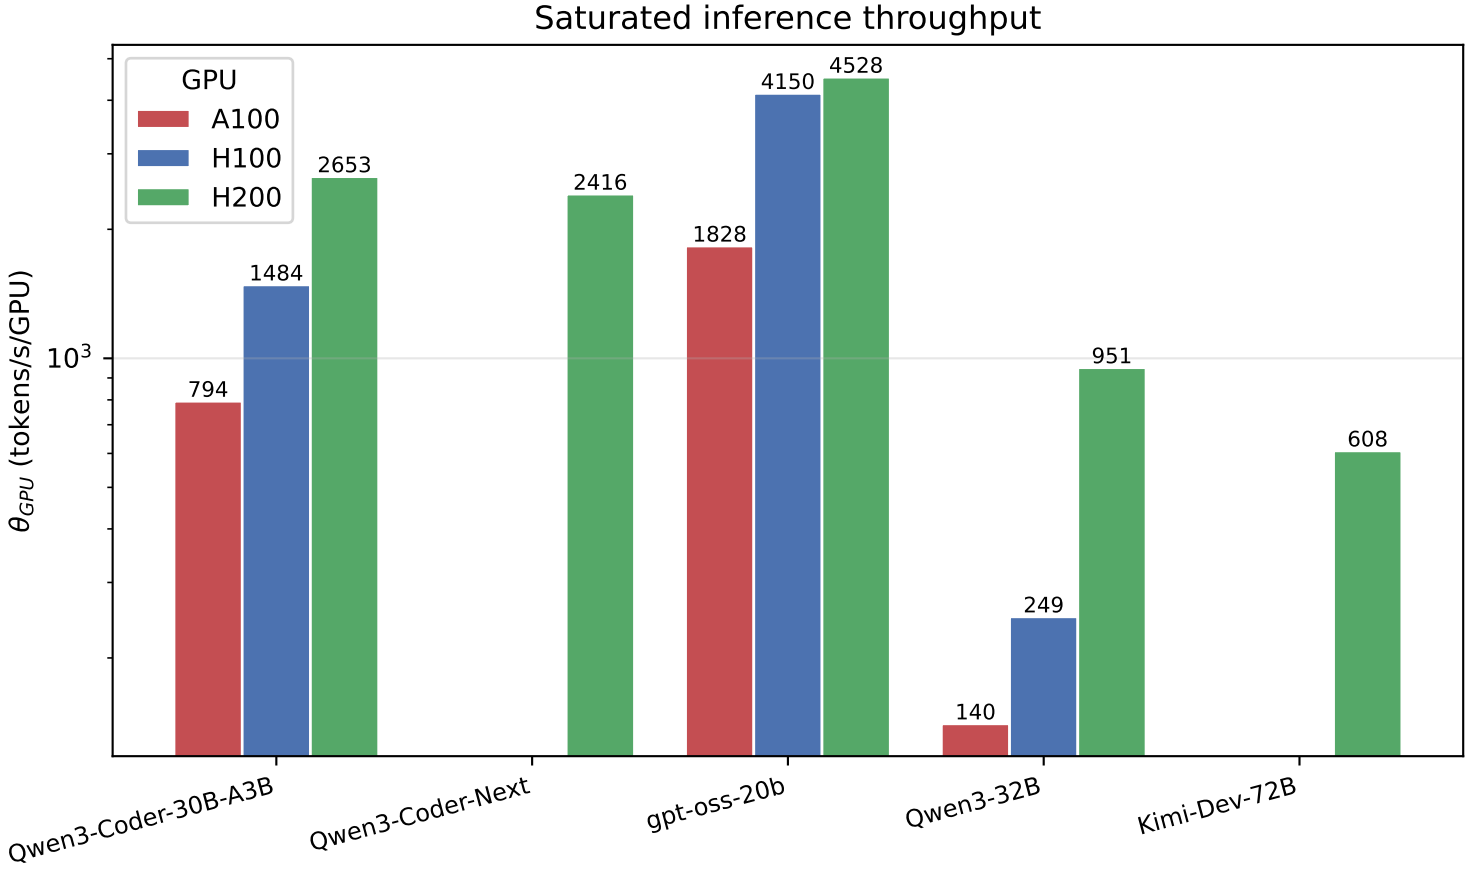
Saturated inference throughput
GPU
A100
H100
H200
$\theta_{GPU}$ (tokens/s/GPU)
$10^3$
794
1484
2653
2416
1828
4150
4528
140
249
951
608
Qwen3-Coder-30B-A3B
Qwen3-Coder-Next
gpt-oss-20b
Qwen3-32B
Kimi-Dev-72B