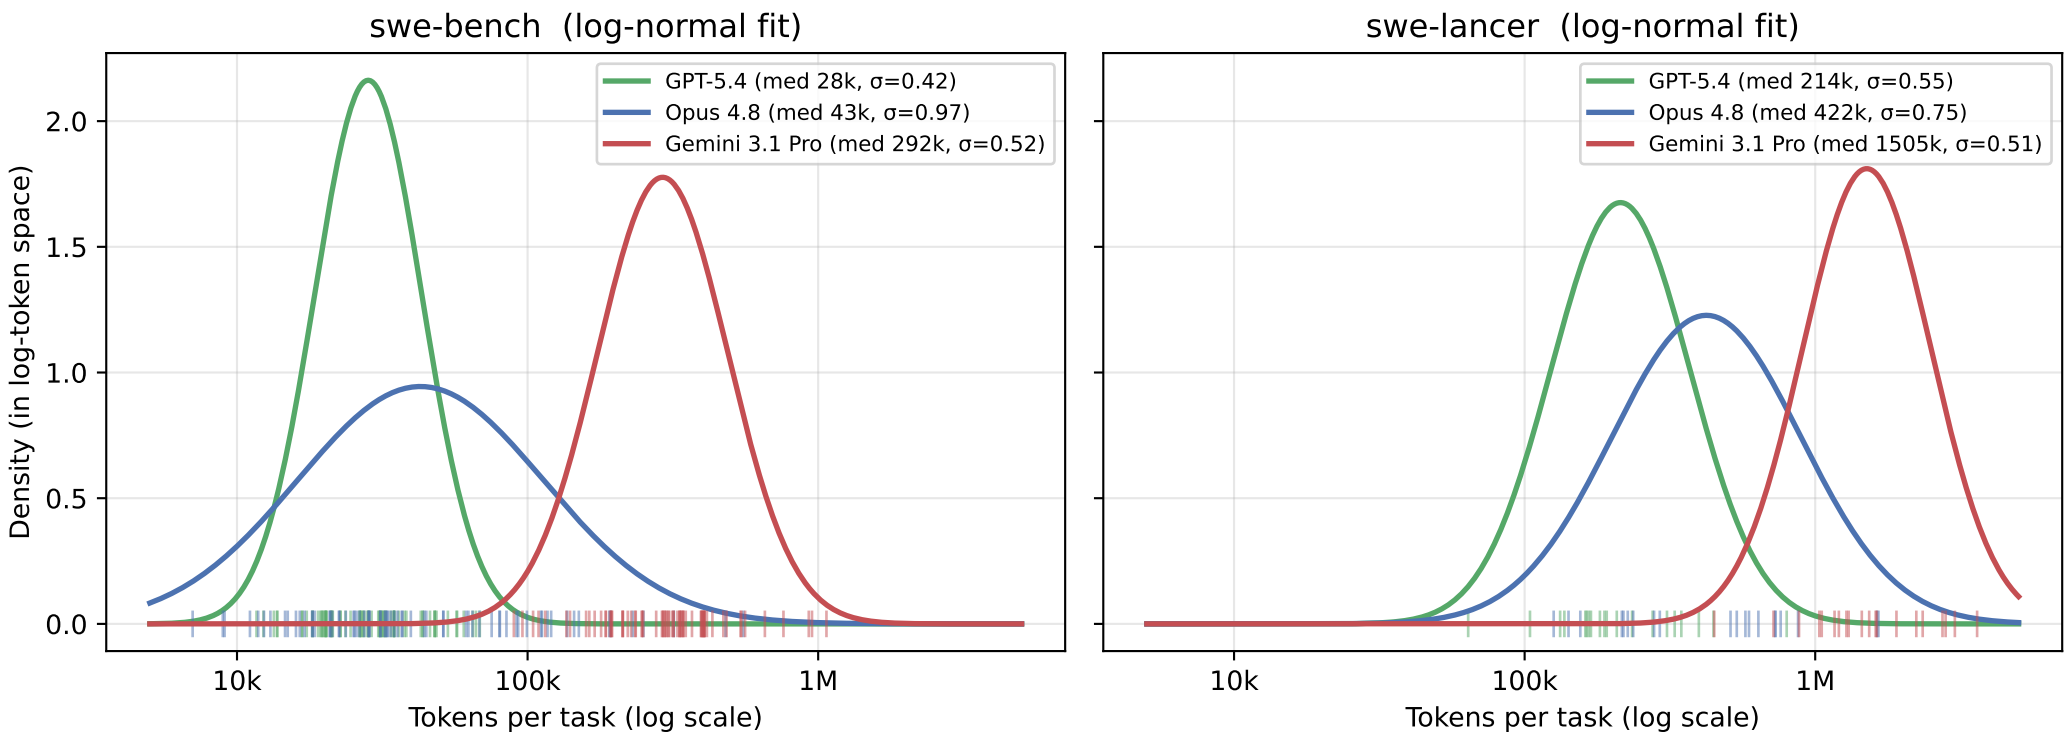
Per-task token-demand distributions, fitted log-normal by model
swe-bench (log-normal fit)
swe-lancer (log-normal fit)
GPT-5.4 (med 28k, σ=0.42)
Opus 4.8 (med 43k, σ=0.97)
Gemini 3.1 Pro (med 292k, σ=0.52)
GPT-5.4 (med 214k, σ=0.55)
Opus 4.8 (med 422k, σ=0.75)
Gemini 3.1 Pro (med 1505k, σ=0.51)
Density (in log-token space)
0.0
0.5
1.0
1.5
2.0
10k
100k
1M
Tokens per task (log scale)