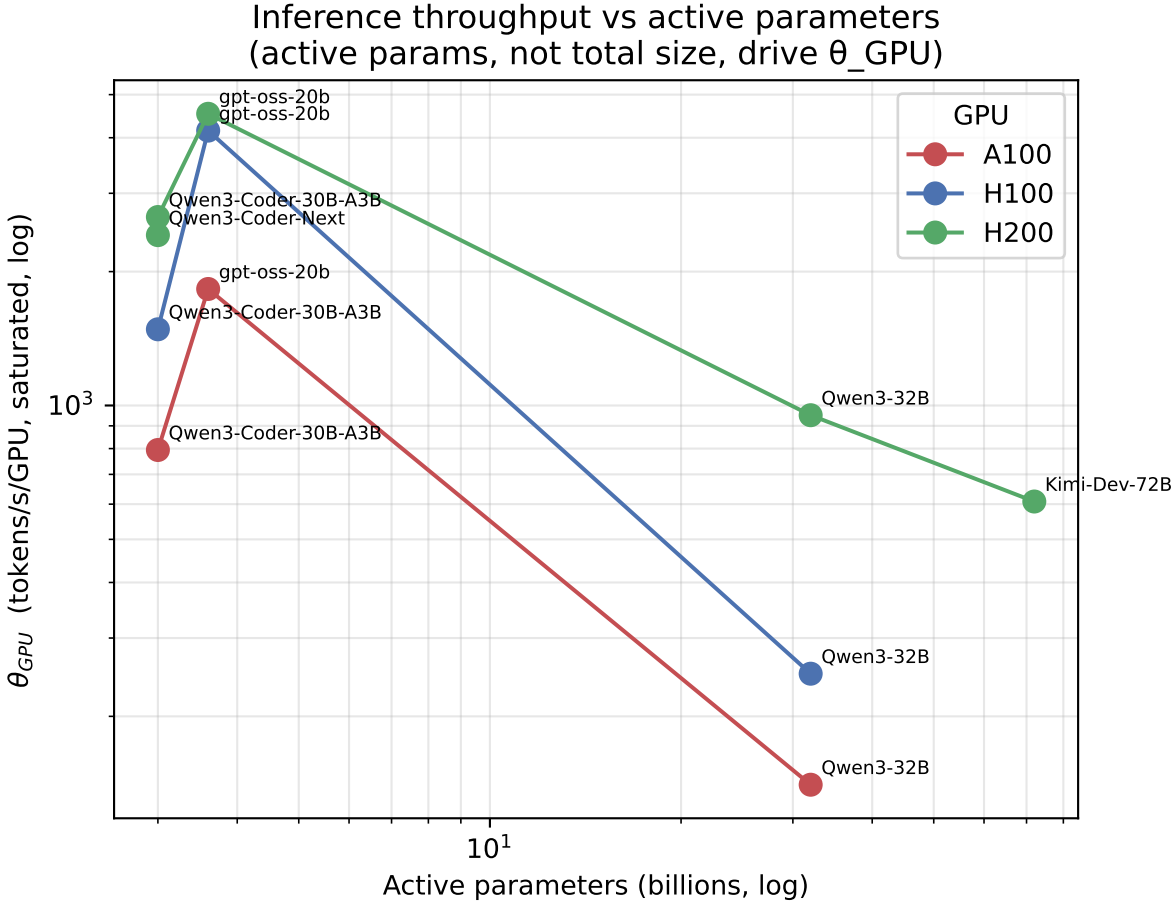
Inference throughput vs active parameters
(active params, not total size, drive θ_GPU)
GPU
A100
H100
H200
gpt-oss-20b
gpt-oss-20b
Qwen3-Coder-30B-A3B
Qwen3-Coder-Next
gpt-oss-20b
Qwen3-Coder-30B-A3B
Qwen3-Coder-30B-A3B
Qwen3-32B
Kimi-Dev-72B
Qwen3-32B
Qwen3-32B
10^3
10^1
θ_{GPU} (tokens/s/GPU, saturated, log)
Active parameters (billions, log)